

FIGURE: Form, Invariance, and Geometric Understanding for Representation Evaluation

Erik J. Bekkers

February 26, 2025

1 Introduction

This document provides details on the FIGURE dataset as hosted on the repository <https://github.com/ebekkers/FIGURE>. Some experiments on this dataset are already reported in its `README.md` file. We propose a synthetic dataset of figures for controlled experiments in robustness evaluation. The dataset consists of point clouds in the Lie group $SE(2)$, with each figure represented as a tree of relative group actions. The primary objective is to assess robustness against texture bias by disentangling geometric pose information from appearance features.

2 Mathematical Formulation

Each figure is modeled as a set of transformations in $SE(2)$, describing the global pose and limb articulations.




Figure 1: Structure of the synthetic stick figure. Each node g_i represents a transformation/pose in $SE(2)$. The root pose g_0 is oriented upward ($\theta_0 = 90^\circ$), and each limb extends hierarchically. Each joint has an associated position x_i , and we denote the end points of each limb with a tilde (\sim), e.g., \tilde{x}_2 denotes the end point of the lower left arm that starts at location x_2 .

2.1 Body Representation

The figure’s global pose is defined as:

$$g_0 = (x_0, \theta_0), \quad x_0 \in \mathbb{R}^2, \quad \theta_0 \in S^1 \equiv SO(2), \quad (1)$$

where x_0 represents the central part of the torso, and θ_0 indicates the figure’s orientation. We use the convention that $\theta = 0$ is aligned with the horizontal axis, and θ rotates counter-clockwise such that $\theta_0 = 90^\circ$ means the figure is standing upright. In this document we provide the angles in degrees ($^\circ$). Each limb’s pose is defined recursively as:

$$g_i = g_{i-1} \cdot \Delta g_i, \quad \Delta g_i \in SE(2) \quad (2)$$

where Δg_i is a relative transformation describing the spatial relation between consecutive limbs.

2.2 Body Joints

Each limb’s pose $g_i = (x_i, \theta_i) \in SE(2)$ consists of a position $x_i \in \mathbb{R}^2$ and an orientation θ_i . We define the projection operator $\pi : SE(2) \rightarrow \mathbb{R}^2$:

$$\pi(g_i) = x_i \quad (3)$$

which extracts the position component from the group element, and essentially distills the *joint* location.

For each limb, we also define an *end-point* \tilde{x}_i , which corresponds to the extremity of the limb segment. Instead of first generating a new group element \tilde{g}_i , we compute the end-point position directly as:

$$\tilde{x}_i = g_i \cdot \Delta \tilde{x}_i \quad (4)$$

where we overload the symbol \cdot to now also represents the group action on \mathbb{R}^2 . The spatial offset $\Delta \tilde{x}_i$ is given by:

$$\Delta \tilde{x}_i = \begin{bmatrix} \ell^{\text{arm}} \\ 0 \end{bmatrix} \quad \text{or} \quad \Delta \tilde{x}_i = \begin{bmatrix} \ell^{\text{leg}} \\ 0 \end{bmatrix} \quad (5)$$

depending on whether i represents an arm or a leg.

2.3 Pose Definitions

Each relative limb transformation Δg_i consists of a translation and a rotation. The pose transformations for different limbs are given in Table. 2.3.

The joint angles θ_i are drawn from probability distributions that include both *natural body variations* and *class-specific variations*, where we define *four classes*: both arms up (up-up), both arms down (down-down), left arm up and right arm down (up-down), and right arm up and left arm down (down-up). See Table. 2.3 for the random variables. Here, $p(\theta_1|\text{up})$ and $p(\theta_3|\text{up-up})$ describe

the angle distributions when the figure belongs to the up-up class, whereas $p(\theta_1|\text{down})$ and $p(\theta_3|\text{down-down})$ describe the angle distributions for the down-down class. The remaining joint distributions capture natural variability in body posture (small angle variations at the joints), independent of class labels.

Table 1: Pose transformations and corresponding relative transformations for each limb.

| Limb | Pose Transformation | Relative Transformation |
|-----------------|------------------------------|---|
| Left upper arm | $g_1 = g_0 \cdot \Delta g_1$ | $\Delta g_1 = (h/2, -w/2, \theta_1)$ |
| Left lower arm | $g_2 = g_1 \cdot \Delta g_2$ | $\Delta g_2 = (\ell^{\text{arm}}, 0, \theta_2)$ |
| Right upper arm | $g_3 = g_0 \cdot \Delta g_3$ | $\Delta g_3 = (h/2, w/2, \theta_3)$ |
| Right lower arm | $g_4 = g_3 \cdot \Delta g_4$ | $\Delta g_4 = (\ell^{\text{arm}}, 0, \theta_4)$ |
| Left upper leg | $g_5 = g_0 \cdot \Delta g_5$ | $\Delta g_5 = (-h/2, -w/2, \theta_5)$ |
| Left lower leg | $g_6 = g_5 \cdot \Delta g_6$ | $\Delta g_6 = (\ell^{\text{leg}}, 0, \theta_6)$ |
| Right upper leg | $g_7 = g_0 \cdot \Delta g_7$ | $\Delta g_7 = (-h/2, w/2, \theta_7)$ |
| Right lower leg | $g_8 = g_7 \cdot \Delta g_8$ | $\Delta g_8 = (\ell^{\text{leg}}, 0, \theta_8)$ |

3 The Shape class and Torch dataloader

The repository contains two important files.

3.1 The Shape class

In `figure_class.py` you'll find a python class `Shape` (based on `numpy`) from which you can sample random figures. This class has the methods `.resample(g, pose_class)` which randomly samples the keypoints according to the distribution of Table. 2.3 given the specified `pose_class` which takes the options `up-up`, `down-down`, `up-down`, `down-up` and applies a group action by g . Setting $g=(0,0,0)$ generates the figure in up-right position at the center of the image.

The class also has the method `.set_pose_from_landmarks(points_r2, resolution)` which sets all the variables from a given set of 13 keypoints. This may be useful in case you want to visualize the result of keypoint detection method, or a point cloud generator.

Then there are 2 visualization methods. `.visualize()` generates a matplotlib figure. `.render(image_ize)` generates a numpy array image of the figure, and also returns the 13 keypoints in pixel coordinates, as well as the 9 $SE(2)$ elements in pixel coordinates. This is the function used to generated the training data provided by the dataloader.

3.2 The Torch dataloader

In `data_loader_torch.py` you'll find the torch dataset `FigureDataset`, which inherits from `torch.utils.data.Dataset`. It returns 4 items: the image, the

Table 2: Random variable distributions for limb orientations and body proportions. Class-specific variations are explicitly indicated.

| Random Variable | = | Distribution |
|-------------------------------------|---|---|
| $p(\theta_1 \mid \text{up-up})$ | = | $U(-85, -10)$ |
| $p(\theta_1 \mid \text{up-down})$ | = | $U(-85, -10)$ |
| $p(\theta_1 \mid \text{down-down})$ | = | $U(-170, -95)$ |
| $p(\theta_1 \mid \text{down-up})$ | = | $U(-170, -95)$ |
| $p(\theta_2)$ | = | $U(-10, 10)$ |
| $p(\theta_3 \mid \text{up-up})$ | = | $U(10, 85)$ |
| $p(\theta_3 \mid \text{up-down})$ | = | $U(95, 170)$ |
| $p(\theta_3 \mid \text{down-down})$ | = | $U(95, 170)$ |
| $p(\theta_3 \mid \text{down-up})$ | = | $U(10, 85)$ |
| $p(\theta_4)$ | = | $U(-10, 10)$ |
| $p(\theta_5)$ | = | $U(170, 225)$ |
| $p(\theta_6)$ | = | $U(-10, 10)$ |
| $p(\theta_7)$ | = | $U(135, 190)$ |
| $p(\theta_8)$ | = | $U(-10, 10)$ |
| $p(w)$ | = | $U(w_{\min}, w_{\max})$ |
| $p(h)$ | = | $U(h_{\min}, h_{\max})$ |
| $p(\ell^{\text{arm}})$ | = | $U(\ell_{\min}^{\text{arm}}, \ell_{\max}^{\text{arm}})$ |
| $p(\ell^{\text{leg}})$ | = | $U(\ell_{\min}^{\text{leg}}, \ell_{\max}^{\text{leg}})$ |

13 keypoints in pixel coordinates (array indices), the 9 $SE(2)$ keypoints, and the class label.

It is also possible to make the shirt color a random variable by setting the parameter `torso_colors`, `color_probabilities`, `color_consistency`. We use this to create biases in the dataset by correlating class labels with a certain color preference (as reflected by the color probabilities). See the `README.md` for more info.